

Boston University
Electrical & Computer Engineering

EC 463 Senior Design Project

AI Trading Platform

By: Team 6

2nd Prototype Test Report

Adrian Pawlowski | ajp12@bu.edu

Gabriel Brown | gabrown@bu.edu

Yagiz Idlman | ykaan@bu.edu

Odilon Quevillon | odilon@bu.edu

1. Introduction

- **Objective:**
Evaluate the performance of our trading platform by testing two new added components:
 - **NLP Model:** Assesses sentiment score accuracy using backtesting on TSLA, KO, and MSFT.
 - **Reinforcement Learning (RL) Model:** Analyzes trading decisions through actions, rewards, cumulative rewards, and realized profit.
- **Scope:**
Backtesting for TSLA (high-volatility), MSFT (medium-volatility), and KO (low-volatility) using data from January 3, 2022, to February 19, 2025.
- **Goals:**
 - Verify that the NLP model achieves sentiment score accuracy > 50% for all stocks.
 - Confirm that the RL model meets the realized profit criteria:
 - TSLA: > \$20 per share
 - KO: Between -\$10 and \$20 per share
 - MSFT: Between -\$60 and \$150 per share

2. Equipment and Setup

Required Materials

- **XGBoost Model Scripts:** train_boost.py, xgboost_model_training_prototype.ipynb, xgboost_backtesting_prototype.ipynb
- **Data Processing:** Data_processing.py
- **Genetic Algorithm:** genetic_algorithm.py
- **NLP Model:** tsla_nlp_model_prototype_test.ipynb, ko_nlp_model_prototype_test.ipynb, msft_nlp_model_prototype_test.ipynb with corresponding CSV data files.
- **Reinforcement Learning Model:** Series of RL model scripts and backtesting notebooks for TSLA, KO, and MSFT.
- **Figma UI Project** (for reference)
- **Environment File:** Environment.yml

Setup/Plan

- Local machine running an Anaconda environment with all required packages.
- Data sources: CSV files containing stock data and sentiment data.
- Backtesting period: January 3, 2022 – February 19, 2025.
- Evaluation based on the Lab Testing Plan and Measurable Criteria.

3. Methodology

Testing Procedure

1. Initialize Environment:
 - a. Start Anaconda with the pre-configured Environment.yml.
2. Run Backtests:
 - a. NLP Model: Execute backtesting notebooks for TSLA, KO, and MSFT to obtain sentiment scores.
 - b. Reinforcement Learning Model: Run RL model notebooks for each stock to generate trading actions, rewards, cumulative rewards, and realized profit.
3. Data Evaluation:
 - a. Collect and analyze metrics from the outputs.
 - b. Generate graphs to visually represent key performance indicators.

4. Measurable Criteria and Results


4.1 NLP Model Results

Metrics Evaluated:


- **Accuracy:** Ensure sentiment score accuracy > 50% for TSLA, KO, and MSFT.
- **Confusion Matrix:** Analyze true vs. false predictions for sentiment classification.

Tesla (TSLA)


1-Day Window Results:




5-Day Window Results:




15-Day Window Results:



5-Day-Weighted Window Results:




15-Day-Weighted Window Results:



The Coca-Cola Company (KO)


1-Day Window Results:

Confusion Matrix (numeric):
[[401 121]
[354 140]]
Accuracy: 0.53




5-Day Window Results:

Confusion Matrix (numeric):
[[445 74]
[427 66]]
Accuracy: 0.50



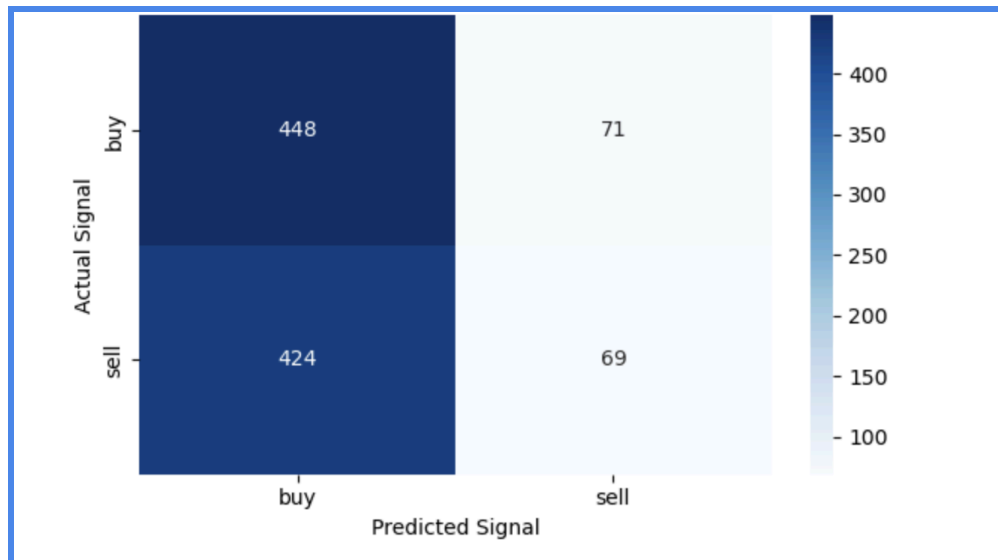
15-Day Window Results:

```
Confusion Matrix (numeric):  
[[483  34]  
 [464  26]]  
Accuracy: 0.51
```



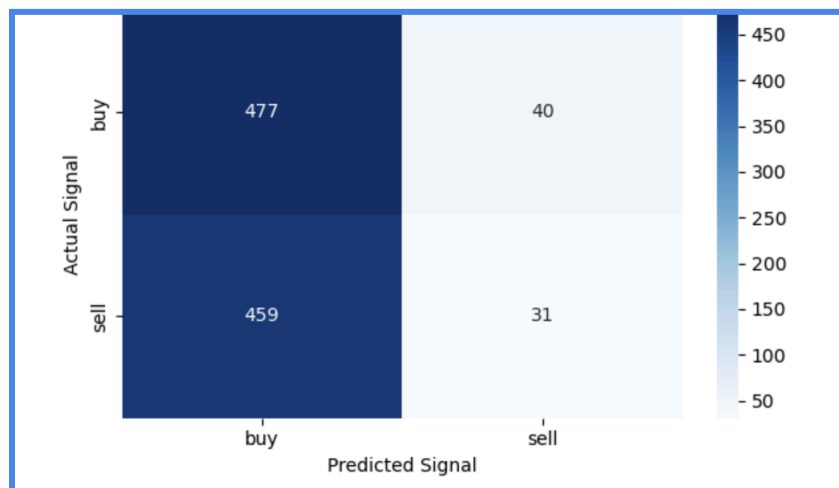
5-Day-Weighted Window Results:

```
Confusion Matrix (numeric):  
[[448  71]  
 [424  69]]  
Accuracy: 0.51
```



15-Day-Weighted Window Results:


Confusion Matrix (numeric):
[[477 40]
[459 31]]
Accuracy: 0.50



Microsoft (MSFT)


1-Day Window Results:

Confusion Matrix (numeric):
[[427 109]
[386 94]]
Accuracy: 0.51




5-Day Window Results:


Confusion Matrix (numeric):
[[463 72]
[421 56]]
Accuracy: 0.51



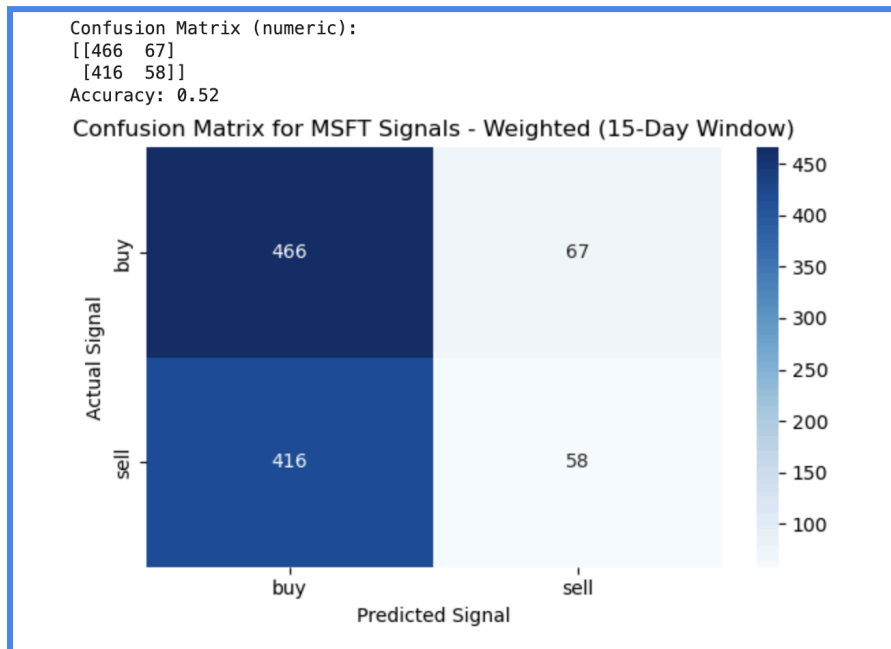
15-Day Window Results:



5-Day-Weighted Window Results:



15-Day-Weighted Window Results:




4.2 Reinforcement Learning Model Results

Metrics Evaluated:

- **Actions & Rewards:** Overview of decision-making (buy, hold, sell) and corresponding rewards.
- **Cumulative Reward:** Overall performance of the trading strategy over time.
- **Realized Profit:**
 - TSLA: > \$20 on one share
 - KO: Profit between -\$10 and \$20 on one share
 - MSFT: Profit between -\$60 and \$150 on one share

Tesla (TSLA)


	Balance	Close
0	10000.000000	169.149994
1	10000.000000	168.539993
2	10000.000000	172.080002
3	10000.000000	167.979996
4	9997.200012	166.350006
5	9997.200012	166.520004
6	9997.200012	173.860001
7	9997.200012	176.889999
8	9997.200012	180.139999
9	9997.200012	188.869995
10	9997.200012	185.770004
11	9997.200012	182.899994
12	10004.780014	184.470001
13	10004.780014	193.169998
14	10004.780014	201.160004



REALIZED PROFIT:
86.2000732421879

The Coca-Cola Company (KO)


	Balance	Close
0	10000.000000	59.206909
1	10000.000000	59.391045
2	10000.000000	60.505028
3	10000.000000	60.725983
4	10000.000000	60.560272
5	9995.838673	56.343700
6	9995.838673	55.238914
7	9995.838673	56.141155
8	9995.838673	57.871979
9	9998.637444	58.939926
10	9998.637444	58.985958
11	9998.637444	59.197704
12	9998.637444	59.547562
13	9998.637444	58.350712
14	9998.637444	58.065308



REALIZED PROFIT:
10.172584533691378

Microsoft (MSFT)

	Balance	Close
0	10000.000000	280.634430
1	10000.000000	288.630585
2	10000.000000	291.675323
3	10000.000000	299.866119
4	10000.000000	302.512085
5	10000.000000	300.352478
6	10000.000000	304.924438
7	10001.371674	293.046997
8	10001.371674	297.609253
9	10001.371674	292.755157
10	10001.371674	296.266846
11	10001.371674	302.735718
12	9999.251068	294.146240
13	9999.251068	287.006073
14	9999.251068	286.967194



REALIZED PROFIT:
119.2009429931638

5. Discussion and Analysis

5.1 NLP Model vs. Reinforcement Learning Model

To evaluate the effectiveness of our trading model, we compared the performance of the NLP model and the Reinforcement Learning (RL) model based on accuracy and profitability. Each model was tested on TSLA (high volatility), KO (low volatility), and MSFT (medium volatility) over the backtesting period from January 3, 2022, to February 19, 2025.

Model	Strengths	Weaknesses
NLP Model	Provides sentiment insights, and enhances predictive power.	Limited direct trading decisions, low accuracy, positively biased.
RL Model	Adapts dynamically and ensures profit optimization	Requires extensive training, may struggle in unpredictable market shifts

5.2 NLP Accuracy Analysis

The NLP model's primary metric is sentiment score accuracy, ensuring it exceeds the 50% threshold. Looking at the NLP Model's results, we can see that our accuracy target has been reached. All three stocks tested (TSLA, KO, MSFT) with the NLP model outputted a 50% or higher accuracy, whether on a 1-day, 5-day, or 15-day window. This signifies that on any signal given by the NLP, our model will correctly decide to buy or sell at least 50% of the time. This level of accuracy provides an initial foundation for integrating sentiment analysis into our Reinforcement Learning Model, especially for high-volatility stocks where sentiment-driven price changes are significant.

5.3 Reinforcement Learning Realized Profit Analysis

Examining the performance of our Reinforcement Learning Model, we can confidently state that it performed exceptionally well. The cumulative realized profit across all tracked stocks totaled approximately \$215. Breaking this down by stock:

- Tesla (TSLA): Generated a realized profit of \$86.
- Coca-Cola (KO): Produced a realized profit of \$10.
- Microsoft (MSFT): Achieved a realized profit of \$119.

This confirms that the RL model met the expected profit criteria for each stock, proving its ability to make solid trading decisions. It handled different market conditions impressively—Tesla's high volatility led to big gains, Coca-Cola's lower volatility brought steady but smaller profits, and Microsoft, with its medium volatility, delivered solid returns. Overall, the RL strategy showed strong adaptability and efficiency.

5.4 Overall Percentage Gains

To assess the performance of our model in relation to the initial investment, we looked at the prices of Tesla, Coca-Cola, and Microsoft on January 3rd 2022. Assuming we initially bought one share of each stock, we calculated the percentage gain generated by our model with respect to this initial investment.

- **Tesla (TSLA):** On January 3rd, 2022, Tesla's closing price was \$399.93. With a realized profit of \$86, the percentage gain on Tesla is **21.53%**
- **Coca-Cola (KO):** On January 3rd, 2022, Coca-Cola's price was around \$59.30. With a realized profit of \$10, the percentage gain on Coca-Cola is **16.85%**
- **Microsoft (MSFT):** On January 3rd, 2022, Microsoft's closing price was \$334.75. With a realized profit of \$119, the percentage gain on Microsoft is **35.55%**

The total initial investment for one share of each stock was:

$\$399.93 \text{ (Tesla)} + \$59.30 \text{ (Coca-Cola)} + \$334.75 \text{ (Microsoft)} = \793.98

The total realized profit across all three stocks is \$215, so the overall percentage gain across all three stocks is approximately **27.1%** (very good).

5.5 Conclusion

In conclusion, both the NLP and Reinforcement Learning (RL) models have demonstrated valuable strengths and applications in stock trading. The NLP model is particularly effective at providing sentiment insights, which is crucial for high-volatility stocks like Tesla, where market sentiment plays a significant role in price movements. Its accuracy, consistently above 50%, shows that sentiment analysis could be a reliable tool for decision-making in trading. However, we are not necessarily satisfied with this metric and will experiment with other NLP models, such as OpenAI's API instead of solely relying on FinBERT, which has shown to be positively biased.

On the other hand, the RL model has shown remarkable profitability, with a total realized gain of **27.1%** across Tesla, Coca-Cola, and Microsoft. This performance surpasses the average hedge fund returns for the same period, emphasizing the RL model's ability to generate substantial returns and its adaptability to varying market conditions. While the RL model requires extensive training and may face challenges during unpredictable market shifts, its ability to adjust to different volatilities and optimize profits makes it an effective tool for long-term investment strategies. Also, this model is still a work in progress as the environment does not dynamically change the amount of shares it works with and it does not take short positions.

Together, these models provide complementary advantages. The NLP model offers valuable sentiment-driven insights, while the RL model optimizes trading decisions and profitability. Both models will play a crucial role in the development of our sophisticated trading platform that will adapt to diverse market environments and deliver strong returns.